

Combining Instance Filtering and Attribute Weighting in Decision Tree and Naïve Bayes Hybrids: An Experimental Study

Ikramul Hasan Moral
Dept. of CSE
United International University
Dhaka, Bangladesh
imoral223489@bscse.uui.ac.bd

Dr. Dewan Md. Farid
Professor, Dept. of CSE & Director, MSCSE Program
United International University
Dhaka, Bangladesh
dewanfarid@cse.uui.ac.bd

Abstract—A published family of hybrid classifiers cleans training data in two ways: a Naïve Bayes model removes training rows that it misclassifies, and a decision tree selects attributes and gives them weights for Naïve Bayes. We rebuilt both cleaning steps and every way of combining them: each step on its own, both steps in either order, and both steps applied to separate copies of the training data and then merged. We evaluated ten public datasets with ten repetitions of stratified 10-fold cross-validation, redid every cleaning step inside each training fold, and scored every held-out row. Our main findings are negative and consistent. Plain C4.5 averaged 86.33% accuracy, and every design that removed instances landed about five points below it, no matter how the steps were arranged. Weighting attributes alone raised Naïve Bayes from 75.93% to 77.35%, the largest gain any cleaning design made over its base classifier, but even that difference is not statistically significant. None of the sixteen paired comparisons, on accuracy or on macro-averaged F1, was significant after Holm correction; the smallest adjusted p -value was 0.44. When we instead cleaned once before cross-validation, the scores of the filter-based designs rose by 13.5 to 16.5 points, because the cleaning step had already seen the test rows and had removed difficult cases from the evaluation. Deletion also fell unequally on the classes: on breast cancer the plain filter removed 50.63% of recurrence-event rows but only 14.51% of the rest. We conclude that combining these two cleaning steps does not improve accuracy under a fair evaluation, and we recommend cleaning inside the training fold, reporting what was deleted, and checking which classes lost training rows.

Index Terms—Naïve Bayes, decision tree, data cleaning, instance filtering, attribute weighting, cross-validation

I. INTRODUCTION

Farid et al. [1] published two hybrid classifiers for multi-class classification. The first cleans training data with Naïve Bayes (NB): it fits NB on the training rows, removes the rows whose predicted label disagrees with the supplied label, and fits a decision tree on the rows that remain. The second cleans the attributes: it fits a decision tree, keeps the attributes that the tree tests, gives a larger weight to attributes used nearer the root, and fits NB on the weighted attributes. A natural question follows: if each step helps on its own, does using both together help more, and does the order matter?

This paper answers that question with a controlled experiment. We rebuilt both published steps from their descriptions,

using the same tools (Weka's C4.5 implementation [2] and a textbook NB), and we tested twelve settings on ten public datasets: the two plain classifiers, each cleaning step on its own with either final classifier, both steps in either order with either final classifier, and a parallel design in which the two steps act on separate copies of the same training fold and the two cleaned copies are merged into one training set. Every setting was evaluated with ten repetitions of stratified 10-fold cross-validation, with all cleaning redone inside each training fold and every held-out row scored.

The experiment produced five findings that we state here and support in Section IV.

- 1) Our rebuild is sound: the plain classifiers land close to the averages of the corresponding 2014 tables, within 2.81 points for C4.5 and 1.40 points for NB.
- 2) Removing instances is the damaging step. All four designs that remove instances and end in a tree score about five points below plain C4.5 (80.89% to 81.21% against 86.33%), whether the removal runs alone, first, second, or in parallel with the attribute step.
- 3) The order of the two steps and the choice to chain them at all do not change the outcome. Swapping the order moves the mean by at most 0.71 points, and chaining is never better than running the two steps separately.
- 4) Weighting attributes alone gives the largest gain over its base classifier: 77.35% for weighted NB against 75.93% for plain NB. The gain does not survive combination, and it is not statistically significant either ($p = 0.131$).
- 5) Cleaning before cross-validation inflates every filter-based score by 13.5 to 16.5 points. The inflated numbers measure which rows the filter removed, not better classification.

A sixth observation concerns what the filters delete: deletion falls unequally on the classes, and removing instances first changes which attributes the tree selects. Section IV-D reports this in detail.

II. THE TWO CLEANING STEPS AND HOW WE COMBINED THEM

A. Removing instances with Naïve Bayes

Given training data D , this step fits NB on D , predicts the label of every training row, and keeps a row only when the prediction agrees with the supplied label. The final classifier, NB or C4.5, is then fitted on the retained rows. One safety rule is in force: if deletion would leave fewer than two classes, the step keeps the original training rows instead. Our runs never triggered this rule.

The rule behind the step is disagreement, not proof of error. A row can be rejected because its label is wrong, but also because the class distributions overlap or because NB does not suit the data. We therefore report which rows were removed rather than calling them noise.

B. Selecting and weighting attributes with a decision tree

This step fits a C4.5 tree on the training data and reads two things from it: which attributes the tree tests, and how near the root each attribute is used. Let d_j be the smallest depth at which the tree tests attribute j , counted with the root at depth one. The weight is

$$w_j = \begin{cases} 1/\sqrt{d_j}, & j \text{ is tested by the tree,} \\ 0, & \text{otherwise.} \end{cases} \quad (1)$$

Attributes with weight zero are dropped. The weighted NB score for class c is

$$s_c(x) = \log P(c) + \sum_j w_j \log P(x_j | c), \quad (2)$$

and NB predicts the class with the largest score. Equation (1) makes attributes tested near the root count more, and Equation (2) applies those weights inside NB.

C. The combined designs

Algorithm 1 shows the two sequential designs. The *instances first* design runs the NB removal step, then fits the selection tree on the retained rows; the *attributes first* design selects and weights attributes on the original training fold, then runs the removal step with weighted NB. In the attributes-first design the same weights (Equation (1)) are used by the filtering model and by the final NB; they are not estimated again after deletion. When the final classifier is C4.5, the weights drop out (a tree does not use them) and only the selected columns carry over.

The parallel design leaves out the chaining. It runs the NB removal step on one copy of the training fold and the selection step on a second copy, both fitted on the original data, and keeps the rows and the columns that the two copies accept. The final classifier is then fitted on that merged set; the final NB uses the tree's weights. This design answers a practical question: if chaining the two steps is supposed to help, the combined design should at least beat doing both separately.

If the selection tree tests no attribute, the design keeps all current columns with unit weights. No such event occurred in

Algorithm 1 A combined cleaning design with final learner L

Require: Training data D , design o , final learner $L \in \{\text{NB}, \text{C4.5}\}$

```

1: if  $o = \text{instances first}$  then
2:    $D \leftarrow \text{REMOVEINSTANCES}(D, \text{plain NB})$ 
3:    $(S, W) \leftarrow \text{SELECTATTRIBUTES}(D, \text{C4.5})$ 
4: else
5:    $(S, W) \leftarrow \text{SELECTATTRIBUTES}(D, \text{C4.5})$ 
6:    $D \leftarrow \text{the columns } S \text{ of } D$ 
7:    $D \leftarrow \text{REMOVEINSTANCES}(D, \text{NB weighted by } W)$ 
8: end if
9: if  $o = \text{instances first}$  then
10:   $D \leftarrow \text{the columns } S \text{ of } D$ 
11: end if
12: Fit  $L$  on  $D$  (NB uses the weights  $W$ )

```

TABLE I
DATASET SIZES IN THE FILES USED FOR EVALUATION

Dataset	Rows	Attributes	Classes
Breast cancer	286	9	2
Contact lenses	24	4	3
Diabetes	768	8	2
Glass	214	9	6
Iris	150	4	3
Soybean	683	35	19
Vote	435	16	2
Image segmentation	2310	19	7
Tic-tac-toe	958	9	2
NSL-KDD	25192	41	22

our runs. At prediction time the test rows keep the selected columns, and every held-out row is scored.

III. EXPERIMENTAL SETUP

A. Data

Table I lists the ten public datasets named in the 2014 study, in the file forms we used. Two of our files differ from the 2014 tables: image segmentation combines two UCI files, and NSL-KDD uses its 20% training file [3]. Glass and NSL-KDD each contain one fewer class than the original study lists. The two files allow comparisons on shared data but not an exact reproduction of the original data preparation.

B. Implementation

C4.5 is Weka J48 3.8.6 with pruning and default settings [2]. NB uses add-one counts for nominal attributes and Gaussian likelihoods for numeric attributes. Attributes stay in their original columns, so tree selection refers to original attributes and not to binary indicator columns. The category vocabularies come from the complete files and stay fixed across folds; a missing nominal value forms its own category. Counts and distributions are estimated on training rows only.

C. The main evaluation: cleaning inside each training fold

The main evaluation uses stratified 10-fold cross-validation, repeated ten times with different random splits, giving 100

folds per dataset and per setting. All twelve settings share the same splits. Inside every training fold, the pipeline refits the NB model that removes rows, the tree that selects attributes, and the final classifier; the held-out fold is then scored with all of its rows. No test row is ever removed, and no learner setting is tuned.

This evaluation is deliberately strict. The cleaning steps see only the training rows of the fold they run in, so they cannot borrow information from the rows they will be scored against.

D. A diagnostic: cleaning once before cross-validation

For comparison, we also ran a weaker procedure: clean the whole dataset once, then cross-validate only the final classifier on the cleaned data. This procedure lets the cleaning step see labels from rows that later end up in test folds, and, for the designs that remove rows, it also shrinks the set of rows that will ever be scored. Its scores are reported only as a diagnostic of what this weaker procedure would exaggerate; they are not performance estimates.

E. Measures and statistical tests

We report accuracy and macro-averaged F1, the unweighted mean of the per-class F1 scores over the classes present in a fold. Each dataset’s score averages its 100 folds, and the summary gives each dataset equal weight; the spread values in the tables are standard deviations over the ten repetition-level averages. For statistical tests we use two-sided Wilcoxon signed-rank tests on the ten per-dataset averages, with exact permutation enumeration. The comparison family covers, for each final classifier, the five cleaning settings against the plain classifier, the two order swaps, and the two sequential-versus-parallel contrasts, sixteen comparisons in total, each computed on accuracy and on macro-F1. Holm’s correction, which rescales p-values for the number of tests, is applied across all thirty-two. We chose this family after seeing the saved results, so the tests are exploratory: a non-significant result does not prove that two settings are equal.

IV. RESULTS AND FINDINGS

A. No cleaning design beats the plain classifiers

Table II shows the main scores and Figure 1 shows each cleaning setting’s change from its base classifier. Plain C4.5 leads on both measures, 86.33% accuracy and 81.92% macro-F1. The four designs that remove instances and end in a tree score between 80.89% and 81.21%, a loss of about five points, and the four versions that end in NB stay within half a point of plain NB (75.42% to 76.39% against 75.93%), with no significant difference among them. The two designs that only touch attributes, with no row removal, stay close to their plain classifiers: selecting attributes for C4.5 costs 0.14 points (86.19% against 86.33%), and weighting attributes for NB gains 1.42 points (77.35% against 75.93%), the largest gain over a base classifier in the whole experiment, though not a significant one ($p = 0.131$).

The per-dataset tables show that the averages hide the worst cases. Table III gives accuracy per dataset. On glass, the

TABLE II
MAIN SCORES (%): MEAN $\pm$ STANDARD DEVIATION OVER THE TEN REPETITION-LEVEL AVERAGES; EACH DATASET HAS EQUAL WEIGHT

Method	Accuracy (%)	Macro-F1 (%)
Plain NB (no cleaning)	75.93 $\pm$ 0.32	69.80 $\pm$ 0.52
Plain C4.5 (no cleaning)	86.33 $\pm$ 0.33	81.92 $\pm$ 0.44
Remove instances only $\rightarrow$ NB	75.42 $\pm$ 0.46	69.08 $\pm$ 0.46
Remove instances only $\rightarrow$ C4.5	81.03 $\pm$ 0.37	75.70 $\pm$ 0.63
Weight attributes only $\rightarrow$ NB	77.35 $\pm$ 0.41	71.73 $\pm$ 0.63
Select attributes only $\rightarrow$ C4.5	86.19 $\pm$ 0.33	81.84 $\pm$ 0.42
Instances first, then attributes $\rightarrow$ NB	75.63 $\pm$ 0.45	69.60 $\pm$ 0.72
Instances first, then attributes $\rightarrow$ C4.5	81.21 $\pm$ 0.29	75.76 $\pm$ 0.58
Attributes first, then instances $\rightarrow$ NB	76.35 $\pm$ 0.34	69.91 $\pm$ 0.60
Attributes first, then instances $\rightarrow$ C4.5	80.89 $\pm$ 0.45	73.93 $\pm$ 0.80
Both steps separately, combined $\rightarrow$ NB	76.39 $\pm$ 0.37	70.47 $\pm$ 0.68
Both steps separately, combined $\rightarrow$ C4.5	81.05 $\pm$ 0.34	75.58 $\pm$ 0.58




Fig. 1. Mean accuracy change from the corresponding plain classifier, over ten datasets with equal weight. Hatched bars end in NB, solid bars in C4.5. Every design that removes instances loses about five points for the tree; nothing else moves the means by more than 1.5 points.

instances-first design with a NB ending falls from 45.95% to 30.93%, and the same design with a C4.5 ending falls from 66.58% to 48.97%. On NSL-KDD the plain tree scores 99.47% and the removal designs score between 87.07% and 89.68%. The single small exception is contact lenses, where the instances-first NB design improves from 73.00% to 83.00%; that dataset has 24 rows, so the result carries little weight. Table IV gives the same picture for macro-averaged F1.

B. Order and combination do not change the outcome

Table V lists the paired dataset-level tests. The largest single difference in the whole family is the attributes-first design’s loss to plain C4.5, -5.44 points with raw $p = 0.014$; after Holm correction across all thirty-two tests it stands at $p = 0.44$. No other comparison comes closer to significance, on accuracy or on macro-F1. Swapping the order moves the NB means by $+0.71$ points ($p = 0.275$) and the C4.5 means by -0.32 points ($p = 0.922$). The sequential designs do not beat doing both separately: all eight sequential-versus-parallel contrasts lie within 0.76 points on accuracy and within 1.65 points on macro-F1, and the only raw p-value below 0.05 among them (-0.76 for instances-first against both-separately

TABLE III

ACCURACY BY DATASET (%). “INSTANCES ONLY” REMOVES TRAINING ROWS WITH PLAIN NB AND THEN FITS THE FINAL CLASSIFIER; “ATTRIBUTES ONLY” SELECTS ATTRIBUTES WITH THE TREE (NB ADDITIONALLY USES THE TREE’S WEIGHTS); “INSTANCES FIRST” AND “ATTRIBUTES FIRST” CHAIN THE TWO STEPS IN THAT ORDER; “BOTH SEPARATELY” APPLIES THE TWO STEPS TO SEPARATE COPIES OF THE TRAINING FOLD AND MERGES THEM.

Dataset	No cleaning	Instances only	Attributes only	Instances first	Attributes first	Both separately
breast-cancer	72.12	72.68	72.43	71.28	72.53	72.46
contact-lenses	73.00	73.00	80.67	83.00	79.33	83.00
diabetes	75.52	74.41	76.73	75.58	75.67	75.81
glass	45.95	37.00	47.84	30.93	37.28	32.29
image-segmentation	79.76	74.92	82.30	74.53	77.32	76.57
iris	95.53	95.33	95.40	95.33	95.87	95.60
nsf-kdd	67.42	75.79	62.80	74.53	75.23	76.16
soybean	90.09	89.00	89.40	85.67	86.66	86.95
tic-tac-toe	69.83	71.89	70.68	70.03	68.07	69.72
vote	90.06	90.13	95.24	95.45	95.49	95.36

Dataset	No cleaning	Instances only	Attributes only	Instances first	Attributes first	Both separately
breast-cancer	74.84	71.77	73.91	71.77	72.08	71.20
contact-lenses	84.67	84.67	84.67	84.67	80.67	84.67
diabetes	74.12	75.12	74.18	75.23	75.57	74.86
glass	66.58	48.83	66.58	48.97	49.15	48.65
image-segmentation	96.81	89.24	96.83	89.26	90.88	89.79
iris	94.27	94.93	94.27	94.93	93.80	94.93
nsf-kdd	99.47	87.43	99.47	89.68	87.07	88.37
soybean	92.71	88.71	92.26	87.89	90.98	88.24
tic-tac-toe	84.74	74.43	84.74	74.55	73.17	74.43
vote	95.08	95.13	95.03	95.13	95.54	95.38

on NB, $p = 0.039$) does not survive the correction. On this family of datasets, where the two steps run makes no measurable difference to accuracy.

C. Cleaning before cross-validation inflates the filter scores

Table VI compares the two procedures for the ten cleaning settings. When the cleaning steps run once on the full dataset before cross-validation, every design that removes rows gains 13.50 to 16.51 points over its score under the main evaluation. The designs that only touch attributes gain 0.72 and 0.33 points. The reason is mechanical rather than informative. Cleaning before the split gives the NB model access to labels from rows that later fall in test folds, and, for the row-removing designs, it deletes rows that the main evaluation would have scored. The larger numbers therefore describe an easier task on a smaller sample. This is the pattern to watch for when a published pipeline reports its cleaning gains without stating where the cleaning ran.

D. Deletion falls unevenly on classes and changes attribute selection

Table VII reports, per dataset, the share of training rows that the two removal designs delete and the share of attributes the selection tree keeps. The removal rate varies widely between datasets: the plain filter deletes 47.31% of glass training rows but 4.14% of iris rows. It also varies between classes within a dataset. On breast cancer, where recurrence events form about 30% of the rows, the plain filter removes 50.63% of recurrence-event training rows against 14.51% of the rest; the weighted filter used in the attributes-first design removes 58.97% and 11.72%. The minority class loses training examples at several times the rate of the majority class, and

on some folds a class disappears from training entirely. The records cannot say whether the removed labels were wrong; they show which rows the filter chose to discard.

The attribute stage reacts to prior deletion. On glass, the tree keeps 97.00% of the attributes when fitted on the original training fold but only 61.00% when fitted after row removal; on vote the share falls from 27.38% to 7.62%. So the two steps are not independent: removing rows first changes what the tree would select. What the interaction does not do is change the accuracy outcome, as Table V shows.

V. PRACTICAL RECOMMENDATIONS

The findings translate into four concrete practices. First, clean inside the training fold. Cleaning before cross-validation inflated every filter-based score by 13.5 to 16.5 points in our experiment, so any reported gain from a cleaning pipeline depends on where the cleaning ran. Second, report what was deleted. A cleaning step that removes rows changes the training population, and per-class deletion records expose effects that mean accuracy hides, such as the 50.63% versus 14.51% class asymmetry on breast cancer. Third, check the classes that lose examples. Rare classes can vanish from training on some folds even with safety rules in place. Fourth, do not expect accuracy from chaining these two cleaning steps. Across ten datasets, both orders, and the parallel design, no arrangement beat the plain classifiers, and the instance-removal step cost about five points wherever it ran.

VI. LIMITATIONS

The benchmark is the 2014 study’s dataset family, which is small and old; contact lenses has 24 rows with classes of four and five rows, and NSL-KDD contains classes with

TABLE IV
MACRO-AVERAGED F1 BY DATASET (%), WITH THE SAME SETTING NAMES AS TABLE III.

Dataset	No cleaning	Instances only	Attributes only	Instances first	Attributes first	Both separately
breast-cancer	64.42	64.22	62.57	61.33	61.57	62.94
contact-lenses	62.31	62.31	72.27	76.47	69.87	76.47
diabetes	72.32	71.43	72.76	71.96	71.05	72.03
glass	44.98	40.72	44.71	36.78	40.42	37.99
image-segmentation	77.78	74.52	81.42	73.40	77.06	75.82
iris	95.48	95.28	95.35	95.29	95.82	95.56
nsl-kdd	34.06	35.87	37.58	35.36	38.18	36.46
soybean	93.13	92.59	91.33	87.27	88.40	89.95
tic-tac-toe	63.83	64.08	64.26	62.93	61.45	62.30
vote	89.66	89.80	95.02	95.25	95.29	95.15

Dataset	No cleaning	Instances only	Attributes only	Instances first	Attributes first	Both separately
breast-cancer	61.81	62.68	60.52	62.65	61.37	61.82
contact-lenses	79.47	79.47	79.47	79.47	72.27	79.47
diabetes	70.71	71.45	70.79	71.56	71.19	70.97
glass	62.00	45.29	62.00	45.45	43.35	45.03
image-segmentation	96.81	88.96	96.83	88.98	90.79	89.53
iris	94.18	94.88	94.18	94.88	93.72	94.88
nsl-kdd	82.25	60.33	82.92	62.05	53.18	60.82
soybean	94.26	92.08	94.10	90.47	92.08	91.21
tic-tac-toe	82.82	66.94	82.82	67.18	66.02	66.94
vote	94.86	94.90	94.81	94.90	95.34	95.17

TABLE V
PAIRED DATASET-LEVEL COMPARISONS. EACH DIFFERENCE IS THE MEAN OVER TEN DATASETS, IN PERCENTAGE POINTS; HOLM p IS CORRECTED ACROSS ALL THIRTY-TWO TESTS. NO COMPARISON IS SIGNIFICANT AT 0.05 AFTER CORRECTION.

Comparison	Acc. diff. (pp)	p	Holm p	F1 diff. (pp)	p	Holm p
Remove instances only vs plain C4.5	-5.30	0.055	1.000	-6.22	0.164	1.000
Select attributes only vs plain C4.5	-0.14	0.562	1.000	-0.07	0.938	1.000
Instances first vs plain C4.5	-5.12	0.055	1.000	-6.16	0.164	1.000
Attributes first vs plain C4.5	-5.44	0.014	0.438	-7.99	0.037	1.000
Both separately vs plain C4.5	-5.28	0.055	1.000	-6.33	0.164	1.000
Attributes first vs instances first (C4.5)	-0.32	0.922	1.000	-1.83	0.232	1.000
Instances first vs both separately (C4.5)	+0.16	0.461	1.000	+0.18	0.461	1.000
Attributes first vs both separately (C4.5)	-0.16	0.770	1.000	-1.65	0.193	1.000
Remove instances only vs plain NB	-0.51	0.652	1.000	-0.71	0.232	1.000
Weight attributes only vs plain NB	+1.42	0.131	1.000	+1.93	0.193	1.000
Instances first vs plain NB	-0.30	1.000	1.000	-0.19	0.557	1.000
Attributes first vs plain NB	+0.42	0.846	1.000	+0.11	0.922	1.000
Both separately vs plain NB	+0.46	0.695	1.000	+0.67	0.846	1.000
Attributes first vs instances first (NB)	+0.71	0.275	1.000	+0.31	0.492	1.000
Instances first vs both separately (NB)	-0.76	0.039	1.000	-0.86	0.055	1.000
Attributes first vs both separately (NB)	-0.05	0.770	1.000	-0.56	0.846	1.000

fewer than ten examples, including one singleton. Some folds therefore lack a class in training or testing, and repetition does not create new examples. Our files differ from the originals in the ways stated in Section III-A, and one discrepancy remains unresolved: our NSL-KDD C4.5 baseline is 99.47% against 71.11% reported in 2014. Removing a row because NB disagrees with its label is not the same as detecting a wrong label, and we have no verified label errors in this data. The nominal category vocabularies come from the complete files, so the estimates do not cover unseen category values. The statistical family was chosen after seeing the results, so the tests are exploratory, and non-significance is not evidence of equality. Finally, the findings hold for these ten datasets, these two classifiers, and these settings; they do not establish how the steps behave elsewhere, nor do they identify which

procedure the 2014 study used.

VII. CONCLUSION

We rebuilt two published data-cleaning hybrids and every way of combining them, including both sequential orders and a parallel design that merges the two steps' separate outputs, and evaluated all twelve settings under an evaluation that refits every cleaning step inside each training fold and scores every held-out row. The combined designs did not beat the plain classifiers on ten public datasets. Removing instances cost about five points of accuracy wherever it ran; weighting attributes alone helped NB by an amount that was not statistically significant and did not survive combination; the order of the steps and the choice to chain them made no measurable difference. Cleaning once before cross-validation

TABLE VI

MEAN ACCURACY (%) UNDER THE MAIN EVALUATION (CLEANING INSIDE EACH TRAINING FOLD) AND UNDER THE DIAGNOSTIC (CLEANING ONCE BEFORE CROSS-VALIDATION). THE DIFFERENCE IS DESCRIPTIVE; THE DIAGNOSTIC CHANGES BOTH THE INFORMATION AVAILABLE AND THE ROWS BEING SCORED.

Design	Final	Inside folds (%)	Before CV (%)	Difference (pp)
Instances only	NB	75.42	91.93	+16.51
Instances only	C4.5	81.03	94.52	+13.50
Attributes only	NB	77.35	78.07	+0.72
Attributes only	C4.5	86.19	86.52	+0.33
Instances first	NB	75.63	90.11	+14.48
Instances first	C4.5	81.21	94.72	+13.51
Attributes first	NB	76.35	91.08	+14.74
Attributes first	C4.5	80.89	96.32	+15.43
Both separately	NB	76.39	91.10	+14.71
Both separately	C4.5	81.05	94.67	+13.62

TABLE VII

CLEANING RECORDS UNDER THE MAIN EVALUATION, AVERAGED OVER 100 FOLDS PER DATASET (%): TRAINING ROWS DELETED BY THE PLAIN FILTER AND BY THE WEIGHTED FILTER, AND ATTRIBUTES KEPT BY THE SELECTION TREE WITHOUT AND WITH PRIOR ROW REMOVAL.

Dataset	Removed, plain filter	Removed, weighted filter	Kept, no prior removal	Kept, after removal
breast-cancer	25.24	25.76	41.56	34.67
contact-lenses	4.28	11.03	66.50	64.00
diabetes	23.67	22.99	84.88	83.12
glass	47.31	46.88	97.00	61.00
image-segmentation	19.88	17.35	65.79	40.16
iris	4.14	4.41	50.75	37.75
nsf-kdd	32.52	37.13	69.32	42.59
soybean	9.18	9.06	52.40	41.89
tic-tac-toe	29.32	28.63	100.00	76.67
vote	9.33	4.46	27.38	7.62

raised the filter-based scores by 13.5 to 16.5 points, an inflation that comes from seeing test rows and deleting difficult cases, not from better classification. The deletion records add the practical warning: filters remove minority-class rows at several times the majority-class rate. For practitioners, the safest use of these cleaning steps is inside the training fold, with the deletions reported and checked per class.

REFERENCES

- [1] D. M. Farid, L. Zhang, C. M. Rahman, M. A. Hossain, and R. Strachan, "Hybrid decision tree and naïve Bayes classifiers for multi-class classification tasks," *Expert Systems with Applications*, vol. 41, no. 4, pp. 1937–1946, 2014.
- [2] I. H. Witten, E. Frank, M. A. Hall, and C. J. Pal, *Data Mining: Practical Machine Learning Tools and Techniques*, 4th ed. Cambridge, MA: Morgan Kaufmann, 2016.
- [3] M. Tavallaei, E. Bagheri, W. Lu, and A. A. Ghorbani, "A detailed analysis of the KDD CUP 99 data set," in *2009 IEEE Symposium on Computational Intelligence for Security and Defense Applications (CISDA)*. IEEE, 2009, pp. 1–6.